

Yuxin Chen

CONTACT INFORMATION

Email: yxxchen@ucdavis.edu
GitHub: <https://github.com/YuxinxinChen>

Website: YuxinxinChen.github.io
Phone: +1-484-790-0788

EDUCATION

UNIVERSITY OF CALIFORNIA, DAVIS, CA, USA

Ph.D., Engineering and Computer Science, October 2016–2023 (Expected)

- Thesis Topic:
Asynchronous Partitioned Global Address Space (PGAS) Style Programming with GPUs
- Adviser: [Professor John D. Owens](#)

KING ABDULLAH UNIVERSITY OF SCIENCE AND TECHNOLOGY, SAUDI ARABIA

M.S., Computer Science, September 2013–2015

- Thesis Topic: *GPU-Accelerated dimension-independent adaptive Metropolis*
- Adviser: [Professor David E. Keys](#)

PROFESSIONAL EXPERIENCE

UNIVERSITY OF CALIFORNIA, DAVIS:

PGAS (Partition Global Addressing Space) with Distributed GPUs 2019–2022

- Led the research and open-source development of novel, **state-of-the-art** asynchronous distributed GPU data structures and algorithms that support irregular applications in areas such as biology, graph analytics, and sparse linear algebra. These solutions leverage Partitioned Global Address Space (PGAS)-style communication between distributed GPUs, allowing for:
 - More efficient communication latency hiding
 - Higher bandwidth utilization
 - Reduced pipeline stalls through increased pipeline scheduling flexibility
- Collaborated with researchers and software engineers to understand their needs and implement efficient solutions for GPU-based distributed systems.
- Created a highly effective programming environment that enabled end-users to efficiently leverage GPUs for irregular communication patterns, contributing to the HPC community's understanding of how to design primitives for accelerator-equipped distributed-memory machines.

GPU utilized Apache Spark 2017–2018

- Developed a GPU version of Spark's fundamental data structure, Resilient Distributed Datasets (RDDs), while adhering to the original design philosophy of Spark. The GPU component was designed to be compatible, efficient, flexible, and principled, enabling seamless integration with existing Spark workflows.
- Developed a highly efficient GPU-accelerated logistic regression, map and reduce operations, etc. Our GPU-accelerated map and reduce operations are twice as fast as its CPU-only Spark version, and the GPU-accelerated logistic regression algorithm outperforms its CPU-only Spark version by a factor of five, improving the overall performance of Spark workflows.
- Collaborated with other members of the team to integrate these enhancements into the Apache Spark codebase, enabling end-users to leverage GPU acceleration to gain deeper insights from their data, improve decision-making, and accelerate time-to-insight.

LAWRENCE LIVERMORE NATIONAL LABORATORY (LLNL):

April–June 2020

- I enhanced the effectiveness of hash table member queries by utilizing fingerprint tables for hash tables, where each fingerprint occupied only 8-16 bits of memory. This approach significantly reduced memory traffic and addressed the Translation Lookaside Buffer (TLB) problem, while ensuring high query accuracy. On RMAT datasets, the solution effectively filtered out false queries with less than 5% false positive rate
- Rather than recreating a larger hash table when an insertion failed, I devised an overflow buffer that could store failed insertions. This approach enabled GPU multi-threading searches of keys, leading to quicker query and insertion times. During testing on RMAT datasets, this approach yielded an overflow ratio of less than 0.1% for varying load factors.
- This project required advanced knowledge of data structures, GPU programming, and algorithm optimization. My contributions to this project resulted in significant performance improvements, demonstrating my ability to develop innovative and efficient solutions to complex problems.

NVIDIA

June–September 2018

- Developed a communication aggregator that facilitates dynamic communication aggregation for gathering, scattering, and randomly accessing data between GPUs, resulting in significant improvements in bandwidth utilization.
- Decoupling communication granularity from computation granularity, the aggregator allows users to select the optimal communication granularity. Additionally, by decoupling the time that generated the message from the time it actually sent the data, the aggregator enables users to code in an algorithm's natural logic flow to send messages of their native size
- The aggregator runs transparently along the application code, minimizing the burden on the user, and allows for sending data both in an eager (immediate) and in a more bandwidth-efficient manner via aggregation.

ADOBE:

June–September 2017

- Implemented a solution for automating the tuning of hyperparameters and system configurations for training deep neural networks. This was achieved by using Bayesian Optimization, which allowed for a more efficient and effective search of the parameter space. By leveraging the auto-tuner for the Inception model on Imagenet, the search space was reduced by about half, leading to faster training of the network. This solution provided a significant improvement over traditional manual hyperparameter tuning methods, which can be time-consuming and less effective in finding optimal settings.

PUBLICATIONS

- [1] **Yuxin Chen**, Benjamin Brock, Serban Porumbescu, Aydın Buluç, Katherine Yelick, and John D. Owens. Atos: A Task-Parallel GPU Scheduler for Graph Analytics. In *Proceedings of the International Conference on Parallel Processing (ICPP 2022)*. Association for Computing Machinery, New York, NY, USA, Article 50, 1–11. doi:10.1145/3545008.3545056.
- [2] **Yuxin Chen**, Benjamin Brock, Serban Porumbescu, Aydın Buluç, Katherine Yelick, and John D. Owens. Scalable irregular parallelism with GPUs: Getting CPUs out of the way. In *Proceedings of the International Conference on High Performance Computing, Networking, Storage and Analysis (SC '22)* IEEE Press, Article 50, 1–16.
- [3] Taylor Groves, Ben Brock, **Yuxin Chen**, Khaled Z. Ibrahim, Lenny Oliker, Nicholas J. Wright, Samuel Williams, and Katherine Yelick. 2020. Performance Trade-offs in GPU Communication: A Study of Host and Device-initiated Approaches. In *2020 IEEE/ACM Performance Modeling, Benchmarking and Simulation of High Performance Computer Systems (PMBS)*. 126–137
- [4] Benjamin Brock, **Yuxin Chen**, Jiakun Yan, John D. Owens, Aydın Buluç, and Katherine Yelick. RDMA vs. RPC for Implementing Distributed Data Structures. In *Proceedings*

of the IEEE/ACM 9th Workshop on Irregular Applications: Architectures and Algorithms, *IA3 2019*, pages 17–22, November 2019.

- [5] **Yuxin Chen**, David Keyes, Kody JH Law, and Hatem Ltaief. "Accelerated dimension-independent adaptive Metropolis." *SIAM Journal on Scientific Computing* 38, no. 5 (2016): S539-S565. doi:10.1137/15M1026432

TALKS

LET THE CHICKENS RULE THEMSELVES.

- The International Conference for High Performance Computing, Networking, Storage, and Analysis (SC22), November 13-18, 2022, Dallas, Texas, USA

ATOS: A TASK-PARALLEL GPU SCHEDULER FOR GRAPH ANALYTICS

- 51st International Conference on Parallel Processing (ICPP22), August 29th to Sept 1st, 2022. Online

TASK-PARALLEL GPU DYNAMIC SCHEDULING FRAMEWORK FOR DYNAMIC IRREGULAR COMPUTATIONS

- Women in High Performance Computing (WHPC), Nov 14, 2021, St. Louis, MO, USA

PERFORMANCE TRADEOFFS IN GPU COMMUNICATION: A STUDY OF HOST AND DEVICE-INITIATED APPROACHES

- 11th IEEE International Workshop on Performance Modeling, Benchmarking and Simulation of High Performance Computer Systems (PMBS20), Nov 9–19, 2020, Virtual

RDMA VS. RPC FOR IMPLEMENTING DISTRIBUTED DATA STRUCTURES

- 9th Workshop on Irregular Applications: Architectures and Algorithms (IA3 2019), November 18, 2019, Denver, CO, USA

TECHNICAL SKILLS

Programming: C, C++, CUDA, Java, Python, Lisp, MATLAB
Operating Systems: Microsoft Windows, Apple OS X, Linux
Other Tools: \LaTeX , UNIX shell scripting, GNU make, CMake, Git

REFERENCES

Professor John D. Owens
Child Family Professor of Engineering and Entrepreneurship
Electrical and Computer Engineering Department of University of California, Davis
E-mail: jowens@ece.ucdavis.edu

Professor Katherine A. Yelick
Robert S. Pepper Distinguished Professor of Electrical Engineering and Computer Sciences
Vice Chancellor for Research at University of California, Berkeley
Senior Faculty Scientist at Lawrence Berkeley National Laboratory
University of California, Berkeley E-mail: yelick@berkeley.edu

Professor Aydın Buluç
Senior Scientist at the Lawrence Berkeley National Lab (Applied Mathematics and Computational Research Division)
Adjunct Assistant Professor of EECS (CS Division) at University of California, Berkeley
E-mail: abuluc@lbl.gov